

Improving CNN Draughts Evaluators using Genetic Algorithms

Student Name: Thien P. Nguyen

Supervisor Name: Stefan Dantchev

Submitted as part of the degree of BSc Computer Science to the

Board of Examiners in the School of Engineering and Computing Sciences, Durham University

January 7, 2018

Abstract —

Background

Presently, competitive Draughts AI players are currently designed to play at a fixed ability. While it has produced very competitive and intelligent players, they require manual modifications in order to improve its performance. This is due to their dependency on pre-defined move databases, where optimal moves are pre-calculated, and recalled when necessary. By combining Neural Networks and Genetic Algorithms, this issue could possibly be solved by creating a player that can grow in ability over time, without the dependency on move-banks.

Aims

The purpose of this project is to explore approaches to tackle the game of English Draughts via the use of machine learning techniques. First, we study previous historical successes in the field, and look at the components that helped build their systems. Then, we look at contemporary methods of computer science that could be used to evolve the historical systems. The project will establish whether this approach provides an effective performance on the game.

Method

The initial population will consist of randomly generated AI players, which will play each other to determine the best player out of the population. The performance of championing AI players at every generation of the genetic algorithm are measured against previous champions. Appropriate algorithms are implemented to detect the overall development of the system's ability to play Checkers.

Proposed Solution.

The proposed solution starts with designing a neural network that evaluates the probability of a particular side winning, given a given state of a checkerboard. This is then used in a algorithm that evaluates future moves to predict the best move at a given position. This, alongside a set of weights for the neural network, creates a player that can evaluate potential moves. Finally, the player is then used on an existing Draughts framework that will provide the player with the ability to play Draughts.

Keywords — AI, Neural Networks, Genetic Algorithms, MiniMax, Alpha Beta Pruning, Draughts

I INTRODUCTION

The intention of this project is to explore the effectiveness of genetic algorithms to improve the evaluation of a neural network. Neural networks are used to determine the performance of two players in a game of checkers. We attempt to use various crossover and mutation strategies to manipulate the neurons of the network in order to increase the accuracy of the measurements. This would allow us to create an effective Draughts playing agent that would have the ability to learn without human input.

Draughts

English Draughts (or Checkers) is a popular 2-player boardgame played on an 8x8 chess-board. Players begin with 12 pieces each, and they are placed on light-coloured squares. Each player takes a turn to move a piece diagonally in one square. They also have the option to capture their opponents piece by moving two consecutive diagonal squares, where the opponents piece is placed immediately opposite the players piece. Pieces can be captured consecutively in a single turn if the moves afford the scenario. In the event that a piece reaches the opposite side of the board from where the piece started with, they are promoted to a King piece. King pieces have the ability to traverse backwards in the same diagonal motion as pawns. A player wins by capturing all of their opponents pieces. A player loses by having all of their pieces captured. A draw occurs when there is both players agree to draw after a three-fold repetition, or a player has pieces on the board but cannot move any of them.

Genetic Algorithms

Genetic algorithms (GAs) are a group of search techniques used to find exact or approximate solutions to optimisation and search problems. It borrows techniques from Charles Darwin's evolutionism theory; individuals are created by the crossover of the genetic information of their parents. Genetic algorithms are a subset of evolutionary algorithms, which is also formed of similar strategies, which also include evolutionary programming [Fogel, 1993][McDonnel, 1993], and genetic programming [Koza, 1991]. Genetic Algorithms focus on the use of genome manipulation via the use of crossover algorithms and mutation methods. Genomes are a metaphor of genetic information, where it typically refers to an 1D array.

Neural Networks

Neural Networks are non-linear statistical data-modelling tools, linking inputs and outputs adaptively in a learning process similar to how the human brain operates. Networks consist of units, described as neurons, joined by a set of rules and weights. The units are defined with characteristics, and appear in layers. The first layer is defined as the input layer, and the last layer being the output. Layers between the two aforementioned are described as hidden layers. Data is analysed by processing them through the layers.

Learning takes place in the form of the manipulation of the weights connecting the units in the layers. This allows it to model complex relationships between inputs and output, and it can also find patterns in the data.

Motivation

Whilst the use of evolutionary algorithms and neural networks have been explored to create draughts players, my intention is to explore the effectiveness of genetic algorithms. My intention is to determine whether it is possible to produce a performant draughts playing agent by the use of GANNs (Genetic Algorithms and Neural Networks).

Deliverables

Minimum

- Implement a CNN
- Implement a Checkers Game Interface
- Implement a genetic algorithm with an evaluation function that consists of a round robin tournament against the population of CNN Evaluators.
- Implement a mini-max algorithm that chooses moves.

Intermediate

- A user-friendly interface to play against the AI
- A monte-carlo search of the move space.
- Analysis of Crossover methods (within Genetic Algorithms)
- Analysis of Mutation methods (within Genetic Algorithms)

Advanced

- Convolutional Neural Network Layer analysis
- The resulting AI can play to an ELO of at least 1200.

Related Work

Arthur Samuel in the 60's pioneered the concept of AI that can play Checkers, through the proposal of using genetic algorithms to produce a Draughts playing agent (Samuel 2000). However, the work described did not consider the use of neural networks (as it was not conceived at the time). Also, it was dependent on a set of heuristics that he devised, which caps the agent's ability to play to the effectiveness of Samuel's heuristics.

The idea of evolving neural networks to play Draughts is based on the success Chellapilla and Fogel had in evolving their own Checkers neural networks using far less sophisticated hardware (Chellapilla & Fogel 1999). Their work, Blondie24, used a neural network as the evaluator function, and used an evolutionary algorithm to evolve the agents. However, new agents are made strictly through the use of mutation. Blondie24's neural network structure consists of a {32, 40, 10, 1} set. Spatial awareness as it takes an immediate input of the positions on the board. This makes it inherently more difficult for the neural network to generate heuristics based on spatial awareness as it is not immediately considered.

II DESIGN

Requirements

Functional Requirements

Code	Description	Priority
F1	A checkers gameboard is created.	High
F2	Agents are able to harness neural networks to assist in their move decision.	High
F3	Offspring agents can be created using parents.	High
F4	The weights and biases of the Agent's neural network are saved to storage.	High
F5	Humans are able to play against the Agents.	High

Non-Functional Requirements

Code	Description
N1	Agents only choose valid, legal moves.
N2	Agents always return a valid, legal move.
N3	Agents are able to play against other agents.

Algorithms and Data Structures

.1 MiniMax Decision Making

In order to choose the best move given a current position, the decision making process for an agent revolves around the use of minimax algorithm. The use of alpha-beta pruning will help to prune unnecessary moves.

The minimax method will have a search depth of 4ply (where the agent will search two moves ahead.) This will allow the agents to have a basic strategy where they can plan their moves in advance. There are inherent tradeoffs with having a higher ply count; where the asymptotic complexity is exponential; it being $O(x^y)$ with x being the branching factor, and y being the depth. The branching factor will consist of the moves from each agent in a given game.

Once the initial system is ready, we can migrate to a hybrid technique that combines Minimax and Monte-Carlo Tree Search (called MCTS-EPT or MCTS with early playout termination) by Lorentz (Lorentz 2016). In his paper, the method is shown to dramatically improve performance; especially in the case for games with especially large branching factors.

.2 Neural Network

In order to evaluate the board, we use a feed-forward multilayer perceptron style neural network. The network contains 4 layers, where the input layer consists of 91 nodes, with the output node having 1. The hidden layers have 40 and 10 nodes respectively. Our input array takes in the form of the grid array of the board. The intention is to weigh the Black pawns with a value of 1, and white pawns as -1.

It is common knowledge that a King piece is worth more than a pawn, but it is disputed about its precise value advantage. For the sake of completeness, a King's

Figure 2: The indexes of the 32 pieces of the input layer are the immediate values of the positions on the board.


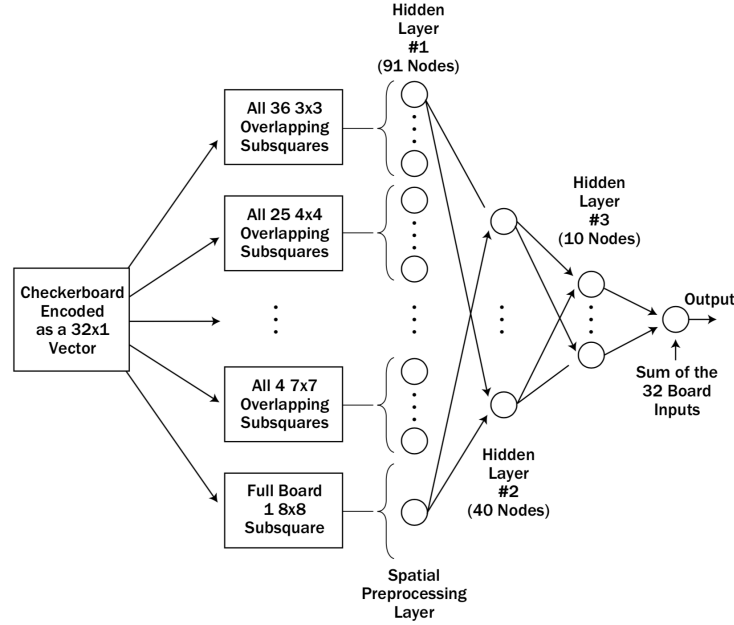


Figure 1: The chosen neural network model. Note that the checkerboard is preprocessed.




piece value is to be weighted at 1.5x the regular pawn value.

To create the input layer, we treat the checkerboard into a 1D array, with the indexes displayed in figure 2. The array is used to calculate all possible subsquares of the checkerboard, ranging from a 3x3 kernel to a 8x8. Each subsquare is summed up to create an input node. There are consequently 91 combinations of the subsquares, thus forming the input layer.

Our initial choice for the activation function is the hyperbolic tangent (a sigmoid), shown in figure 3. As an activation function, it suffices but there are inherent issues, discussed by Hinton.

.3 Genetic Algorithms

Figure 3: graph of sigmoidal function $f(x) = \frac{1}{1+e^{-5x}}$



The genetic algorithm is the heart of the learning strategy of the system. Here we discuss the various algorithms that form the collection of GA strategies.

.4 Population Generation

The initial population will consist of randomly generated weights and biases of the neural network, with values from -1,1 inclusive. For a population size of 15, the next generation is created using the best five agents from the current generation. They will con-

tinue to play in the next generation. The agents are also chosen as a base to create 10 new agents from.

The next eight players are generated through the use of crossover strategies. The weights of the 1st and 2nd place agents are used as input to the crossover strategy and will generate 4 offsprings. Two are reciprocal crossover representations from the crossover, and the other two being directly mutated from the parents themselves. Another four children will be created using the same strategy, with the 2nd and 3rd agent's weights. The remaining two will be direct mutations of the 4th and 5th place agents.

.5 Tournament Selection

The tournament selection process revolves around each agent in the population playing 5 games as Black, against randomly selected opponents. Each game lasts a maximum of 100 moves; if a winner is not deduced from the game after each player places 100 moves, then a draw is called. Draws are also called when there is a five-fold move repetition from both players. A win is worth 2 points, a draw being none and a loss being -1 points. Both the agent and its opponent receives a score. Scores are tallied up at the end of the game. Players are sorted by the number of points they scored. The best players will have the highest number of points.

.6 Coefficient Mutation

Each weight of the neural network will be incremented by a random value that is created using the following formula, where $WeightP$ is the current weight, and K represents the number of weights and biases in the neural network:

$$WeightN = WeightP + \frac{1}{\sqrt{2 * \sqrt{K}}}$$

The weights, as explained earlier will have a hard cap of [-1, 1]. This would consequently mean that the mutation is not controlled, and dependent on the number of weights in the system; The more weights in the network implies a less significant mutation.

.7 Crossover Strategy

Two offsprings are created per a pair of parents, with each offspring being the reciprocal crossover of each other. The weights of both parents (now each treated as a 1D array of coefficients), are divided contingent on the number of weights and biases for a given layer. Each layer should be treated separately to reduce the potential dependency on a purely randomly generated neural network. For each set of weights in a given layer, the following algorithm represents the crossover process:

Choice of Programming Language

There are several contenders, with each having their inherent benefits. C++ is a notable choice due to it's relatively lower level architecture, support for popular machine learning packages (most notably Google's TensorFlow.) In terms of performance, C++ trumps most languages. C++

Algorithm 1 Crossover Strategy

```
1: procedure CROSSOVER
2:    $stringlen \leftarrow \text{length of } string$ 
3:    $i \leftarrow patlen$ 
4: top:
5:   if  $i > stringlen$  then return false
6:    $j \leftarrow patlen$ 
7: loop:
8:   if  $string(i) = path(j)$  then
9:      $j \leftarrow j - 1$ .
10:     $i \leftarrow i - 1$ .
11:    goto loop.
12:   close;
13:    $i \leftarrow i + \max(delta_1(string(i)), delta_2(j))$ .
14:   goto top.
```

has notable parallelised packages, (A popular library is OpenMP) which can assist in the overall performance of the system. However, it would be difficult to write, due to my unfamiliarity with the language. Also, programs written in this language are less portable. It is not suitable for running on university machines without the use of a sandbox. Improperly handled bugs can cause a fatal error on university machines.

Javascript is a contender; Node.JS is a very powerful and popular package manager, *npm*. It is difficult to write multi-processed programs as Node.JS runs on a single thread by nature. It also lacks the support of popular machine learning libraries and performs relatively slower in some programming operations.

I have chosen Python 3.6 due to my familiarity, and the support of popular scientific packages including NumPy and other machine learning tools. Python is also portable with a very wide compatibility; for instance it is pre-installed on all popular UNIX machines and also has support from the university machines. Python development will be on Visual Studio Code, which again is a familiar tool and is also suited to the project.

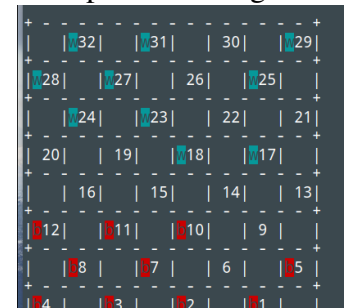
Object Oriented approaches are taken for the majority of the components of the system, ranging from the neural network library to the tournament system. Data structures are implemented using their own classes and methods where applicable. The modularity of object oriented programming provides the affordance of easier debugging and testing.

Players weights (for their neural networks) are stored in two forms, one of which is to be stored on an MongoDB NoSQL instance, and another local copy in JSON. This allows the individual agents to be played against humans.

Interface

As our intention is to find an agent that plays Draughts, having a relatively friendly user interface is not necessarily important, i.e. a simple text-based interface will suffice. Due to the inherently computational strain the project requires, an interface based around having a command-line interface (CLI) the system to be

Figure 4: An example CLI interface demonstrating a game of draughts between user input and an Agent.



quickly up and running, without the need for package dependencies, configurations and set-ups. It also reduces the amount of processing needed to render the most relevant statistics related to the project. The simulation would show information such as estimated finishing times, current generation count, scores of players in a given generation and the cumulative score of progress of the system.

When it comes to humans playing against the agent, the checkerboard can also be rendered using ASCII plaintext, users can make inputs through text (in the console or terminal); where the game will show human users the possible moves that a person can take. This system will be used to play against the agent, where the human will take inputs the moves on behalf of the agent's on-line opponent. A possible rendition is shown in figure 4.

Tools

Initial runs will operate on a 1-ply load in order to determine the stability of the system on a 4-core Intel i5 6200u with 12GB's of memory. Development and debugging will occur on this machine. Once testing has proven to be stable and lacking in errors, the system would run on Durham's MIRA (128-core Intel) distributed system with a 6-ply heavy load. In order to keep simulations running on MIRA, MOSH is used to maintain a consistent connection to MIRA. The end champion is then transferred to the initial machine in order to be played against by human input.

Testing and Evaluation

The combined GANN (Genetic Algorithm/Neural Networks) are evaluated under competition style conditions. At each generation, each agent plays 5 games, with their opponent being from randomly chosen from the generation pool. Point scores are measured by 2,0,-1 where 2 is a win, 0 is a draw, and -1 is a loss. There is a hard cap of 50 moves, where the game is considered a draw if after the game hasn't ended after 50 moves from each player.

At the end of a given generation, we measure growth of performance using the champion of the generation. Presently we will use the mean of means approach. When a new champion is generated, it is played against the previous 5 champions from earlier generations. 6 games are played for each previous champion, with 3 being as Black, and 3 being White. A mean score is calculated from those 6 games. The overall performance of the current champion is the mean of the 5 sets of games. A positive improvement is when the mean of means are greater than 0.

Point Score for the champion games are measured by 1,0,-1 where a Win counts as 1 point and -1 for a loss. The weights are scaled differently to the regular tournament in order to accurately portray the difference between previous champions.

At the end of the generation run, the end player will be used to compete against human players on various online multiplayer checkers websites in order to determine an accurate ELO rating of the system.

References

- Chellapilla, K. & Fogel, D. (1999), 'Evolving Neural Networks to Play Checkers without Expert Knowledge', *IEEE Transactions on Neural Networks* **10**(6), 1382–1391.
URL: <http://ieeexplore.ieee.org/abstract/document/809083>
- Lorentz, R. (2016), 'Using evaluation functions in Monte-Carlo Tree Search', *Theoretical Computer Science* **644**, 106–113.
URL: <http://linkinghub.elsevier.com/retrieve/pii/S0304397516302717>
- Samuel, A. L. (2000), 'Some studies in machine learning using the game of checkers', *IBM Journal of Research and Development* **44**(1.2), 206–226.