

Playing Draughts using Neural Networks and Genetic Algorithms

Thien Nguyen

Department of Computer Science
Durham University

January 16, 2018

Outline

Problem Description

Motivation

Related Work

Approach

Curent Progress

Conclusion

Problem Description

A problem in Computer Science

Presently, competitive Draughts AI players are currently designed to play at a fixed ability.

While it has produced very competitive and intelligent players, they require manual modifications in order to improve its performance.

This is due to their dependency on pre-defined move databases, where optimal moves are pre-calculated, and recalled when necessary.

By combining Neural Networks and Genetic Algorithms, this issue could possibly be solved by creating a player that can grow in ability over time, without the dependency on move-banks.

Motivation

Why have I chosen to tackle this?

- Enjoyed the AI Search module
- Want to learn about Machine Learning (unfortunately not an option this year)
- I love board games!

Related Work

Similar works of art but no cigar

Samuel (59')

Uses Genetic Algorithms to improve coefficients of a set of heuristics to evaluate Draughts games.

Blondie24 (97')

Uses an Evolutionary Algorithm and Neural Networks to evaluate Draughts games. (Quite similar!)

Giraffe (15')

Uses contemporary machine learning techniques to train a Neural Network to evaluate Chess games.

Current Approaches

How will I tackle this?

- ▶ Evaluate Checkerboard
- ▶ Choosing the Best Move
- ▶ Generating Agents
- ▶ Choosing the best Agents
- ▶ Making better Agents
- ▶ ...
- ▶ Profit!

Evaluating Checkerboards





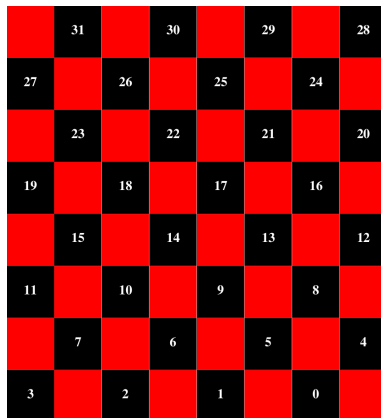
Figure: An example state of a checkerboard.

Neural Networks



$$Output = ActivationFunction((Input * weight) + Bias)$$


Checkerboard



	31		30		29		28
27		26		25		24	
	23		22		21		20
19		18		17		16	
	15		14		13		12
11		10		9		8	
	7		6		5		4
3		2		1		0	

Figure: The indexes of the 32 pieces of the input layer are the immediate values of the positions on the board.

Neural Networks







Figure: graph of tanh function $f(x) = \frac{2}{1+e^{-x}} - 1$

Choosing moves



Monte Carlo

- ▶ Choose a random move at a node
- ▶ Choose another random move
- ▶ Keep choosing random moves n amount of times
- ▶ after n moves, evaluate!
- ▶ (Traditionally you'd randomly choose to the end)

The intuition behind this is that it tends to lean towards better moves probabilistically.

Generating Agents

An agent is a generated set of weights for the neural network.

Tournament

1. Generate a population of random agents
2. Make agents play each other
3. Order agents by the amount of points scored
4. The best few agents are chosen to stay on for the next tournament
5. make new agents from those best few
6. the losers are destroyed
7. repeat step 2-6 with the new agents until satisfied

Crossover Mechanism

Two offsprings would be created from a pair of parents, with each offspring being the reciprocal crossover of each other. The weights of both parents (now each treated as a 1D array of coefficients), are divided contingent on the number of weights and biases for a given layer. Each layer should be treated separately to reduce the potential dependency on a purely randomly generated neural network. For each set of weights in a given layer, the following algorithm represents the crossover process:

Mutation

Weight and biases of an agent's neural network will increment by a random value that is created using the following formula, where $WeightP$ is the current weight, K represents the number of weights and biases in the neural network, and m representing a random floating point in the range of $[-1,1]$:

$$WeightN = WeightP + \frac{m}{\sqrt{2 * \sqrt{K}}}$$

The weights, as explained earlier will have a soft maximum of $[-1, 1]$. This would consequently mean that the mutation is not controlled, and dependent on the number of weights in the system; The more weights in the network implies a less significant mutation.

Current Progress

What have I done already?

- ▶ An agent has been created and trained
- ▶ It plays relatively well (anecdotally)
- ▶ Not necessarily effective at end game performances...

Remaining Work

What do I still need to do?

- ▶ Measure the effectiveness of genetic algorithms
- ▶ Optimise training (enforce threefold repetition)
- ▶ Train the system for a bit longer
- ▶ Measure the system's performance against other players

References

Nanos gigantum humeris insidentes

References!